

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import accuracy_score, confusion_matrix,
precision_score, recall_score, f1_score, auc, precision_recall_curve
from sklearn import preprocessing

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, BatchNormalization
from tensorflow.keras import optimizers
```

```
In [3]: data=pd.read_csv('Churn_Modelling.csv')
data.head(10)
```

Out[3]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMembe
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	
5	6	15574012	Chu	645	Spain	Male	44	8	113755.78	2	1	
6	7	15592531	Bartlett	822	France	Male	50	7	0.00	2	1	
7	8	15656148	Obinna	376	Germany	Female	29	4	115046.74	4	1	
8	9	15792365	He	501	France	Male	44	4	142051.07	2	0	
9	10	15592389	H?	684	France	Male	27	2	134603.88	1	1	

```
In [4]: data.tail(10)
```

```
Out[4]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive
9990	9991	15798964	Nkemakonam	714	Germany	Male	33	3	35016.60	1	1	
9991	9992	15769959	Ajuluchukwu	597	France	Female	53	4	88381.21	1	1	
9992	9993	15657105	Chukwualuka	726	Spain	Male	36	2	0.00	1	1	
9993	9994	15569266	Rahman	644	France	Male	28	7	155060.41	1	1	
9994	9995	15719294	Wood	800	France	Female	29	2	0.00	2	0	
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2	1	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	

```
In [5]: data=data.drop(['RowNumber','CustomerId','Surname'],axis=1)  
data.head()
```

```
Out[5]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

```
In [6]: data.shape
```

```
Out[6]: (10000, 11)
```

```
In [7]: data.isnull().sum()
```

```
Out[7]: CreditScore      0  
        Geography      0  
        Gender         0  
        Age            0  
        Tenure         0  
        Balance        0  
        NumOfProducts  0  
        HasCrCard      0  
        IsActiveMember 0  
        EstimatedSalary 0  
        Exited         0  
        dtype: int64
```

```
In [8]: data.isna().sum()
```

```
Out[8]: CreditScore      0  
        Geography      0  
        Gender         0  
        Age            0  
        Tenure         0  
        Balance        0  
        NumOfProducts  0  
        HasCrCard      0  
        IsActiveMember 0  
        EstimatedSalary 0  
        Exited         0  
        dtype: int64
```

```
In [9]: data.duplicated().sum()
```

```
Out[9]: 0
```

In [10]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   CreditScore         10000 non-null  int64  
 1   Geography           10000 non-null  object  
 2   Gender              10000 non-null  object  
 3   Age                 10000 non-null  int64  
 4   Tenure              10000 non-null  int64  
 5   Balance             10000 non-null  float64 
 6   NumOfProducts       10000 non-null  int64  
 7   HasCrCard           10000 non-null  int64  
 8   IsActiveMember      10000 non-null  int64  
 9   EstimatedSalary     10000 non-null  float64 
10   Exited              10000 non-null  int64  
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB
```

In [11]: data.describe()

Out[11]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881	0.203700
std	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818	0.402760
min	350.000000	18.000000	0.000000	0.000000	1.000000	0.000000	0.000000	11.580000	0.000000
25%	584.000000	32.000000	3.000000	0.000000	1.000000	0.000000	0.000000	51002.110000	0.000000
50%	652.000000	37.000000	5.000000	97198.540000	1.000000	1.000000	1.000000	100193.915000	0.000000
75%	718.000000	44.000000	7.000000	127644.240000	2.000000	1.000000	1.000000	149388.247500	0.000000
max	850.000000	92.000000	10.000000	250898.090000	4.000000	1.000000	1.000000	199992.480000	1.000000

```
In [12]: data['Exited'].value_counts(normalize=True)
```

```
Out[12]: 0    0.7963  
         1    0.2037  
         Name: Exited, dtype: float64
```

```
In [13]: data['Geography'].value_counts(normalize=True)
```

```
Out[13]: France    0.5014  
         Germany   0.2509  
         Spain     0.2477  
         Name: Geography, dtype: float64
```

```
In [14]: data[data['Geography']=='France']['Exited'].value_counts(normalize=True)
```

```
Out[14]: 0    0.838452  
         1    0.161548  
         Name: Exited, dtype: float64
```

```
In [15]: data[data['Geography']=='Germany']['Exited'].value_counts(normalize=True)
```

```
Out[15]: 0    0.675568  
         1    0.324432  
         Name: Exited, dtype: float64
```

```
In [16]: data[data['Geography']=='Spain']['Exited'].value_counts(normalize=True)
```

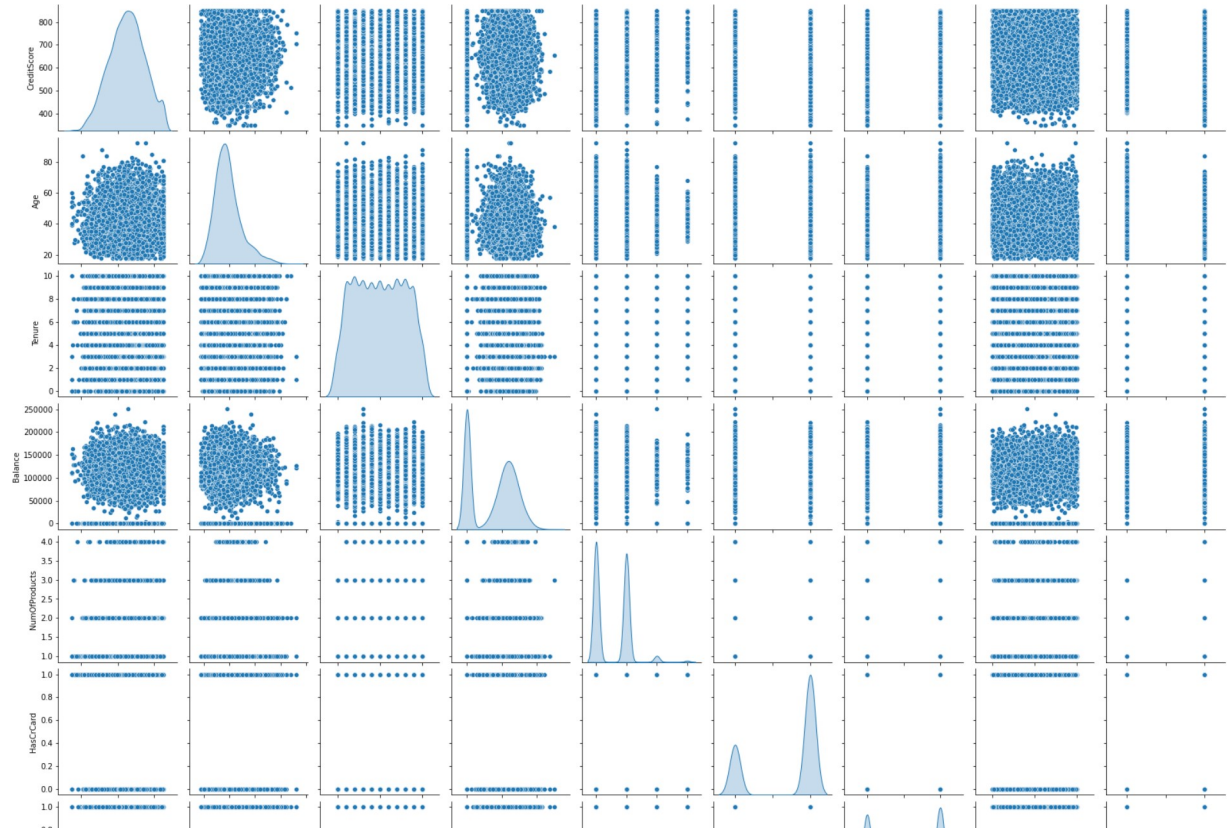
```
Out[16]: 0    0.833266  
         1    0.166734  
         Name: Exited, dtype: float64
```

```
In [17]: data['Gender'].value_counts(normalize=True)
```

```
Out[17]: Male      0.5457  
         Female    0.4543  
         Name: Gender, dtype: float64
```


```
In [18]: sns.pairplot(data,diag_kind='kde')
```

```
Out[18]: <seaborn.axisgrid.PairGrid at 0x1d8aa4368e0>
```



```
In [19]: plt.figure(figsize = (8,8))
sns.heatmap(data.corr(),annot=True)
```

```
Out[19]: <AxesSubplot:~>
```




```
In [20]: data1 = pd.get_dummies(data, columns=['Geography', 'Gender'])
data1.head()
```


```
Out[20]:
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited	Geography_France	Geography
0	619	42	2	0.00	1	1	1	101348.88	1	1	
1	608	41	1	83807.86	1	0	1	112542.58	0	0	
2	502	42	8	159660.80	3	1	0	113931.57	1	1	
3	699	39	1	0.00	2	0	0	93826.63	0	1	
4	850	43	2	125510.82	1	1	1	79084.10	0	0	

```
In [21]: data1.dtypes
```

```
Out[21]: CreditScore      int64
Age                    int64
Tenure                 int64
Balance                float64
NumOfProducts          int64
HasCrCard              int64
IsActiveMember         int64
EstimatedSalary        float64
Exited                 int64
Geography_France        uint8
Geography_Germany       uint8
Geography_Spain         uint8
Gender_Female           uint8
Gender_Male             uint8
dtype: object
```

```
In [22]: for i in data1.columns:
sns.barplot(x='Exited',y=i,data=data1)
plt.show()
```



```
In [23]: x=data1.drop('Exited',axis=1)
x.head()
```

Out[23]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Geography_France	Geography_Germany
0	619	42	2	0.00	1	1	1	101348.88	1	
1	608	41	1	83807.86	1	0	1	112542.58	0	
2	502	42	8	159660.80	3	1	0	113931.57	1	
3	699	39	1	0.00	2	0	0	93826.63	1	
4	850	43	2	125510.82	1	1	1	79084.10	0	

```
In [24]: x.dtypes
```

```
Out[24]: CreditScore      int64
Age                    int64
Tenure                int64
Balance              float64
NumOfProducts        int64
HasCrCard             int64
IsActiveMember        int64
EstimatedSalary      float64
Geography_France      uint8
Geography_Germany     uint8
Geography_Spain       uint8
Gender_Female         uint8
Gender_Male           uint8
dtype: object
```

```
In [25]: x=x.astype('float64')
x.dtypes
```

```
Out[25]: CreditScore      float64
Age                    float64
Tenure                float64
Balance              float64
NumOfProducts        float64
HasCrCard             float64
IsActiveMember        float64
EstimatedSalary      float64
Geography_France      float64
Geography_Germany     float64
Geography_Spain       float64
Gender_Female         float64
Gender_Male           float64
dtype: object
```

```
In [26]: y=data1['Exited']  
y.head()
```

```
Out[26]: 0    1  
1    0  
2    1  
3    0  
4    0  
Name: Exited, dtype: int64
```

```
In [27]: # Split the data up in train and test sets  
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.30, random_state=75)  
print(x_train.shape)  
print(x_test.shape)  
print(y_train.shape)  
print(y_test.shape)  
  
(7000, 13)  
(3000, 13)  
(7000,)  
(3000,)
```

```
In [28]: from sklearn.preprocessing import StandardScaler  
  
# Define the scaler  
scaler = StandardScaler().fit(x_train)  
  
# Scale the train set  
x_train = scaler.transform(x_train)  
  
# Scale the test set  
x_test = scaler.transform(x_test)
```

```
In [39]: model = Sequential()

# Adding the layres with 13 inputs and fully connected neurons.
# using 'ELU' and 'ReLU' activation functions in the hidden layers
# using one sigmoid function at the output as it's a classification model
model.add(Dense(128, input_shape = (13,), activation = 'elu'))
model.add(Dense(64, activation = 'relu'))
model.add(Dense(32, activation = 'swish'))
model.add(Dense(16, activation = 'elu'))
model.add(Dense(64, activation = 'relu'))
model.add(Dense(32, activation = 'swish'))
model.add(Dense(16, activation = 'elu'))
model.add(Dense(8, activation = 'relu'))
model.add(Dense(4, activation = 'swish'))
model.add(Dense(1, activation = 'sigmoid'))

# choose the optimizer, Learning rate, Loss function,metrics
optm=optimizers.Adam(lr=0.0015)
model.compile(optimizer = optm, loss = 'binary_crossentropy', metrics=['accuracy'])
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
dense_10 (Dense)	(None, 128)	1792
dense_11 (Dense)	(None, 64)	8256
dense_12 (Dense)	(None, 32)	2080
dense_13 (Dense)	(None, 16)	528
dense_14 (Dense)	(None, 64)	1088
dense_15 (Dense)	(None, 32)	2080
dense_16 (Dense)	(None, 16)	528

dense_17 (Dense)	(None, 8)	136
dense_18 (Dense)	(None, 4)	36
dense_19 (Dense)	(None, 1)	5

=====

Total params: 16,529

Trainable params: 16,529

Non-trainable params: 0

C:\Users\DRISTI\anaconda3\lib\site-packages\keras\optimizers\optimizer_v2\adam.py:114: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.

super().__init__(name, **kwargs)

```
In [40]: epoch=850
model_hist=model.fit(x_train, y_train,batch_size = 1000,validation_split = 0.2, epochs=epoch, verbose = 1)
model_hist
hist = pd.DataFrame(model_hist.history)
hist['epoch'] = model_hist.epoch
print(hist)
plt.plot(hist['accuracy'])
plt.plot(hist['val_accuracy'])
plt.plot(hist['val_loss'])
plt.plot(hist['loss'])
plt.legend(("train" , "valid","val_loss","loss") , loc =0)
```

```
Epoch 1/850
6/6 [=====] - 2s 41ms/step - loss: 0.6551 - accuracy: 0.7729 - val_loss: 0.6015 - val_ac
curacy: 0.7879
Epoch 2/850
6/6 [=====] - 0s 13ms/step - loss: 0.5674 - accuracy: 0.7907 - val_loss: 0.5187 - val_ac
curacy: 0.7879
Epoch 3/850
6/6 [=====] - 0s 10ms/step - loss: 0.4955 - accuracy: 0.7907 - val_loss: 0.4838 - val_ac
curacy: 0.7879
Epoch 4/850
6/6 [=====] - 0s 10ms/step - loss: 0.4783 - accuracy: 0.7907 - val_loss: 0.4558 - val_ac
curacy: 0.7879
Epoch 5/850
6/6 [=====] - 0s 13ms/step - loss: 0.4546 - accuracy: 0.7907 - val_loss: 0.4381 - val_ac
curacy: 0.7879
Epoch 6/850
6/6 [=====] - 0s 14ms/step - loss: 0.4407 - accuracy: 0.7907 - val_loss: 0.4216 - val_ac
curacy: 0.7879
Epoch 7/850
6/6 [=====] - 0s 14ms/step - loss: 0.4301 - accuracy: 0.7907 - val_loss: 0.4127 - val_ac
```

```
In [31]: y_preds = model.predict(x_test)

# Predict the results using 0.5 as a threshold
print((y_preds>0.5).astype(int))

# Identify the model evaluation loss and accuracy
results = model.evaluate(x_test, y_test)

94/94 [=====] - 0s 1ms/step
[[0]
 [0]
 [0]
 ...
 [0]
 [0]
 [0]]
94/94 [=====] - 0s 2ms/step - loss: 2.4743 - accuracy: 0.8193
```